


Data flow diagram

Tuesday, September 15, 2026 2:49 PM



Assets: Exec-policy configuration, API key, agent context, VM files
 ① ② ③ ④

Principals: User, agent, webpage
 ① ② ③

Trust boundaries: ①, ②, ③, ④

Threats/Controls:

- ① Indirect prompt injection via webpage — Exec-policy requires human approval regardless of what the model is told
- ② Credential leakage — Key kept only in the shim env. variable, nowhere else
- ③ Argument injection into safe_marker.sh — The script accepts the one hardcoded argument
- ④ AI agent reads/writes outside the intended directory — Script has the output path hardcoded
- ⑤ Model executes unauthorized tool call — the allowlist in the exec-policy blocks unapproved calls
- ⑥ Unauthorized access to agent via network — NAT only networking, no external inbound access